

**Read the book two ways: layer-by-layer (parts) or question-by-question (traceability)
each chapter supplies a derived quantity, which feeds a decision**

Part I — The Token

↓ per-token price reduction

↓ what to characterize

Part II – The Workload

capacitance
↓ which is a measure of gate

Part III – The System

↓ fit / complete vs. embry-bound

↓ GPUs ↓ **swirling stack**

Part IV — The Architecture

↓ synthesis deployment / reject

↓ fleet saturation/cloud/API

Part V — The Fleet

$H = \frac{1}{\rho} \left(\frac{1}{\rho} \right)^{\frac{1}{\rho}} \left(\frac{1}{\rho} \right)^{\frac{1}{\rho}} \left(\frac{1}{\rho} \right)^{\frac{1}{\rho}}$

↓ fleet size ↓ multi-depth routing

file → **in2** + rollback
↓ QPS vs. host rollout

Part VI – The Architect

both ways → red/green probe
↓ where to open the problem design

ACM w/ 2026 strategy