

01 / Source



Text

user prompt
characters

02 / Tokenize

Tokenizer

BPE splits
corpus-driven

split
raw string

emit
integer IDs

Token IDs

integers from vocab
ordered

lookup
dense vector

03 / Embed

Embedding

ID → d-dim vector
one vec/token

attend
queries past K, V

04 / Attend

Attention

Q · K mix V
vs past

append K, V
remember prefix

05 / Remember

KV Cache

past K, V stored
per layer × head