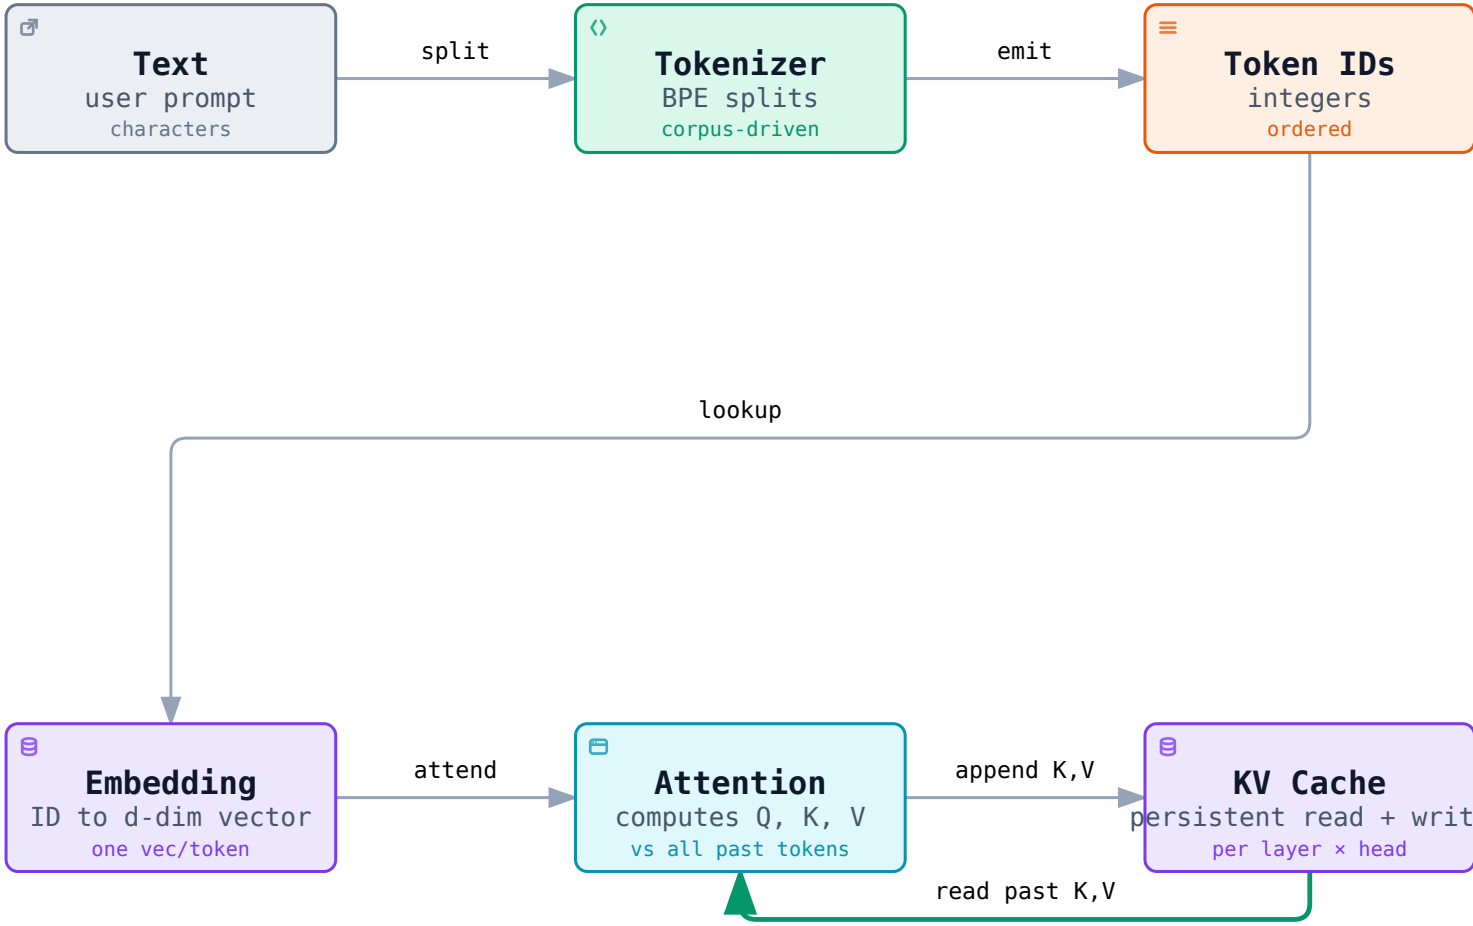
Text
user prompt
characters
split
Tokenizer
BPE splits
corpus-driven
emit
Token IDs
integers
ordered
lookup
Embedding
ID to d-dim vector
one vec/token
attend
Attention
computes Q, K, V
vs all past tokens
append K,V
KV Cache
persistent read + write
per layer × head
read past K,V