


Dense 70B



ALL 70B params active

every token

weight residency $\approx$ 140 GB
KV grows linearly with context

same KV cache (attention is still dense)

MoE (8 experts)



KV growth identical to dense

weight residency $\approx$ 28 GB
KV grows identically to dense

same KV cache (attention is still dense)