Dense 70B

ALL active

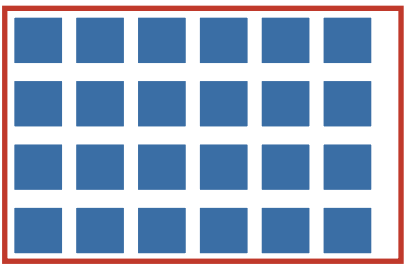

resident ≈ 140 GB

active per token ≈ 140 GB

every token activates all 70B params

KV grows linearly with context

KV: identical (dense = MoE)

controlled: same attention geometry, context length, and cache dtype

MoE (8 experts)

top-2 active

resident ≈ full expert set

active per token ≈ 28 GB

only top-2 of 8 experts compute per token

KV grows linearly with context

KV: identical (dense = MoE)

controlled: same attention geometry, context length, and cache dtype