


## Dense 70B



**ALL 70B params active**

every token

weight residency  $\approx$  140 GB

KV grows linearly with context

same KV cache (attention is still dense)

## MoE (8 experts)



KV growth identical to dense

weight residency  $\approx$  28 GB

KV grows identically to dense

same KV cache (attention is still dense)