


Dense 70B

ALL active



■ resident ≈ 140 GB

■ active per token ≈ 140 GB

every token activates all 70B params

KV grows linearly with context

e (matched att

MoE (8 experts)

top-2 active



■ resident $\approx$ full expert set

■ active per token ≈ 28 GB

only top-2 of 8 experts compute per token

KV grows linearly with context

e (matched att