


Dense 70B



ALL 70B params active

every token

weight residency $\approx$ 140 GB

KV grows linearly with context

same KV cache (attention is still dense)

MoE (8 experts)



top-2 active $\rightarrow$ $\sim$ 14B

KV growth identical to dense

weight residency $\approx$ 28 GB

KV grows identically to dense

same KV cache (attention is still dense)