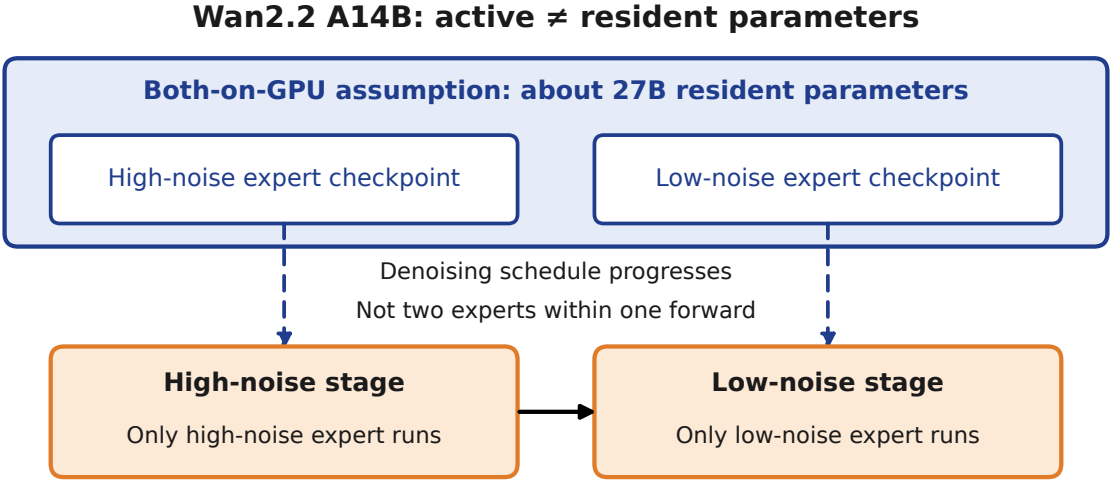
Wan2.2 A14B: active ≠ resident parameters
Both-on-GPU assumption: about 27B resident parameters
High-noise expert checkpoint
Low-noise expert checkpoint
Denoising schedule progresses
Not two experts within one forward
High-noise stage
Only high-noise expert runs
Low-noise stage
Only low-noise expert runs
About 14B active per forward evaluation.
If both experts stay on GPU: about 27B resident.
Schematic: stage widths do not encode step counts.
Offload lowers GPU residency but adds weight transfers.