

KV cache size vs context, by attention geometry (full-MHA vs GQA)

