


KV cache size vs context, by attention geometry (full-MHA vs GQA)



—●— full-MHA FP16 (bound, ~2.62 MB/tok)

—■— full-MHA FP8 (~1.42 MB/tok, 54% of BF16)

—▲— GQA FP16 (~0.33 MB/tok)

—▼— GQA FP8 (~0.18 MB/tok)

..... 8xH100 raw (640 GB)

--- KV budget ~436 GB (640 – 140 w – 64 rt)