


KV cache size vs context, by attention geometry (full-MHA vs GQA)



—●— full-MHA FP16 (bound, ~ 2.62 MB/tok)

—■— full-MHA FP8 (~ 1.42 MB/tok, 54% of BF16)

—▲— GQA FP16 (~ 0.33 MB/tok)

—▼— GQA FP8 (~ 0.18 MB/tok)

..... 8xH100 raw (640 GB)

— — KV budget ~ 436 GB (640 — 140 w — 64 rt)