

# Memory floor: inference vs fine-tuning (70B)

full fine-tune exceeds an 8×H100 host:  
needs multi-node or offload

