


Memory Tetris: how the 8×H100 host (640 GB) fills with context

640 GB pool = 8×H100 HBM



9.2K context **32K context** **128K context**
baseline (weights + runtime) is context-independent, the KV cache is the lever that grows with context. FP16 ~2.5 MB/token [2° DERIVED].

- KV cache (grows w/ context)
- weights 140 GB
- runtime / NCCL ~64 GB