


Memory Tetris: how the 8×H100 host (640 GB) fills with context

640 GB pool = 8×H100 HBM



baseline (weights + runtime) is context-independent; the KV cache is the lever that grows with context. FP16 ~2.5 MB/token [2° DERIVED]

-  kv cache (grows w/ context)
-  weights 140 GB
-  runtime / NCCL ~64 GB