

Expert parallel (all-to-all)

GPU A
experts 1-5

GPU B
experts 6-10

GPU C
experts 11-15

EP composes with the per-stage mesh

the per-stage DP $\times$ TP mesh

the same GPUs also take part in EP all-to-all composition

Pipeline (DP $\times$ TP $\times$ PP)

PP stage 1
a DP $\times$ TP mesh

PP stage 2
a DP $\times$ TP mesh

PP stage 3
a DP $\times$ TP mesh