

Expert parallel (all-to-all)



GPU A
experts 1-5



GPU B
experts 6-10



GPU C
experts 11-15

EP composes with the per-stage mesh



the per-stage DP × TP mesh
the same GPUs also take part in EP all-to-all
composition

Pipeline (DP × TP × PP)



PP stage 1
a DP × TP mesh



PP stage 2
a DP × TP mesh



PP stage 3
a DP × TP mesh