

<>

Single GPU

one accelerator · local memory

outgrows 1 GPU

forces local

<>

Single host

GPU rack node · NVLink fabric

KV overflows floor

privacy

<>

Multi-host

several nodes · cluster fabric

fleet demand

cost pressure

<>

Cluster / fleet

many nodes · network spine