


FlashAttention: the same math, far less HBM traffic



same $Q \cdot K^T \cdot V$ math and same result — the gain is fewer HBM round-trips, not cheaper attention