




- self-host fleet ( $\sim \$148\text{K/mo}$ )
- cloud on-demand ( $\sim \$72\text{K/mo}$ )
- managed API ( $\$20.80/1\text{K}$ , linear)

self-hosted flat (reserved,  $\sim 20$ -host fleet); managed API linear;

cloud on-demand @ 40% duty, 5-host avg bursting to  $\sim 20$  — excludes staff (see §16.4).