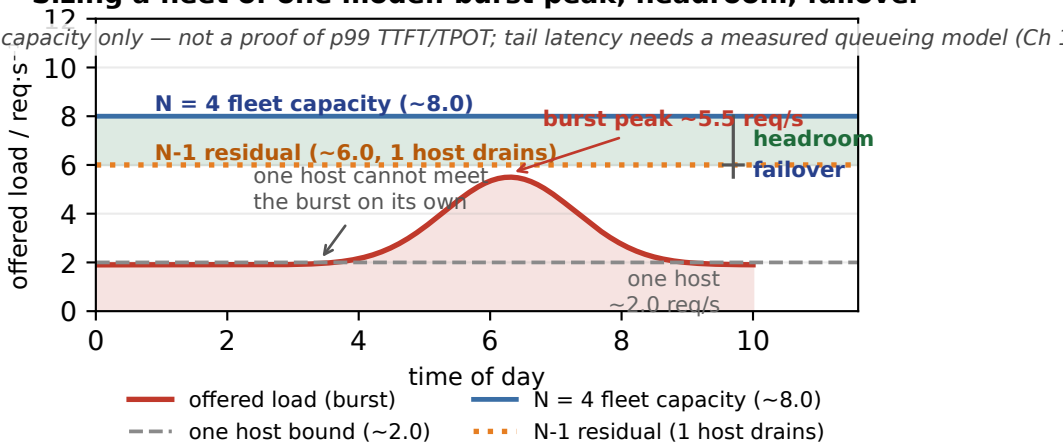
Sizing a fleet of one model: burst peak, headroom, failover
fered-load capacity only — not a proof of p99 TTFT/TPOT; tail latency needs a measured queueing model (Ch 1
offered load / req·s⁻¹
12
10
8
6
4
2
0
N = 4 fleet capacity (~8.0)
burst peak ~5.5 req/s
headroom
N-1 residual (~6.0, 1 host drains)
failover
one host cannot meet
the burst on its own
one host
~2.0 req/s
0
2
4
6
8
10
time of day
offered load (burst)
N = 4 fleet capacity (~8.0)
one host bound (~2.0)
N-1 residual (1 host drains)