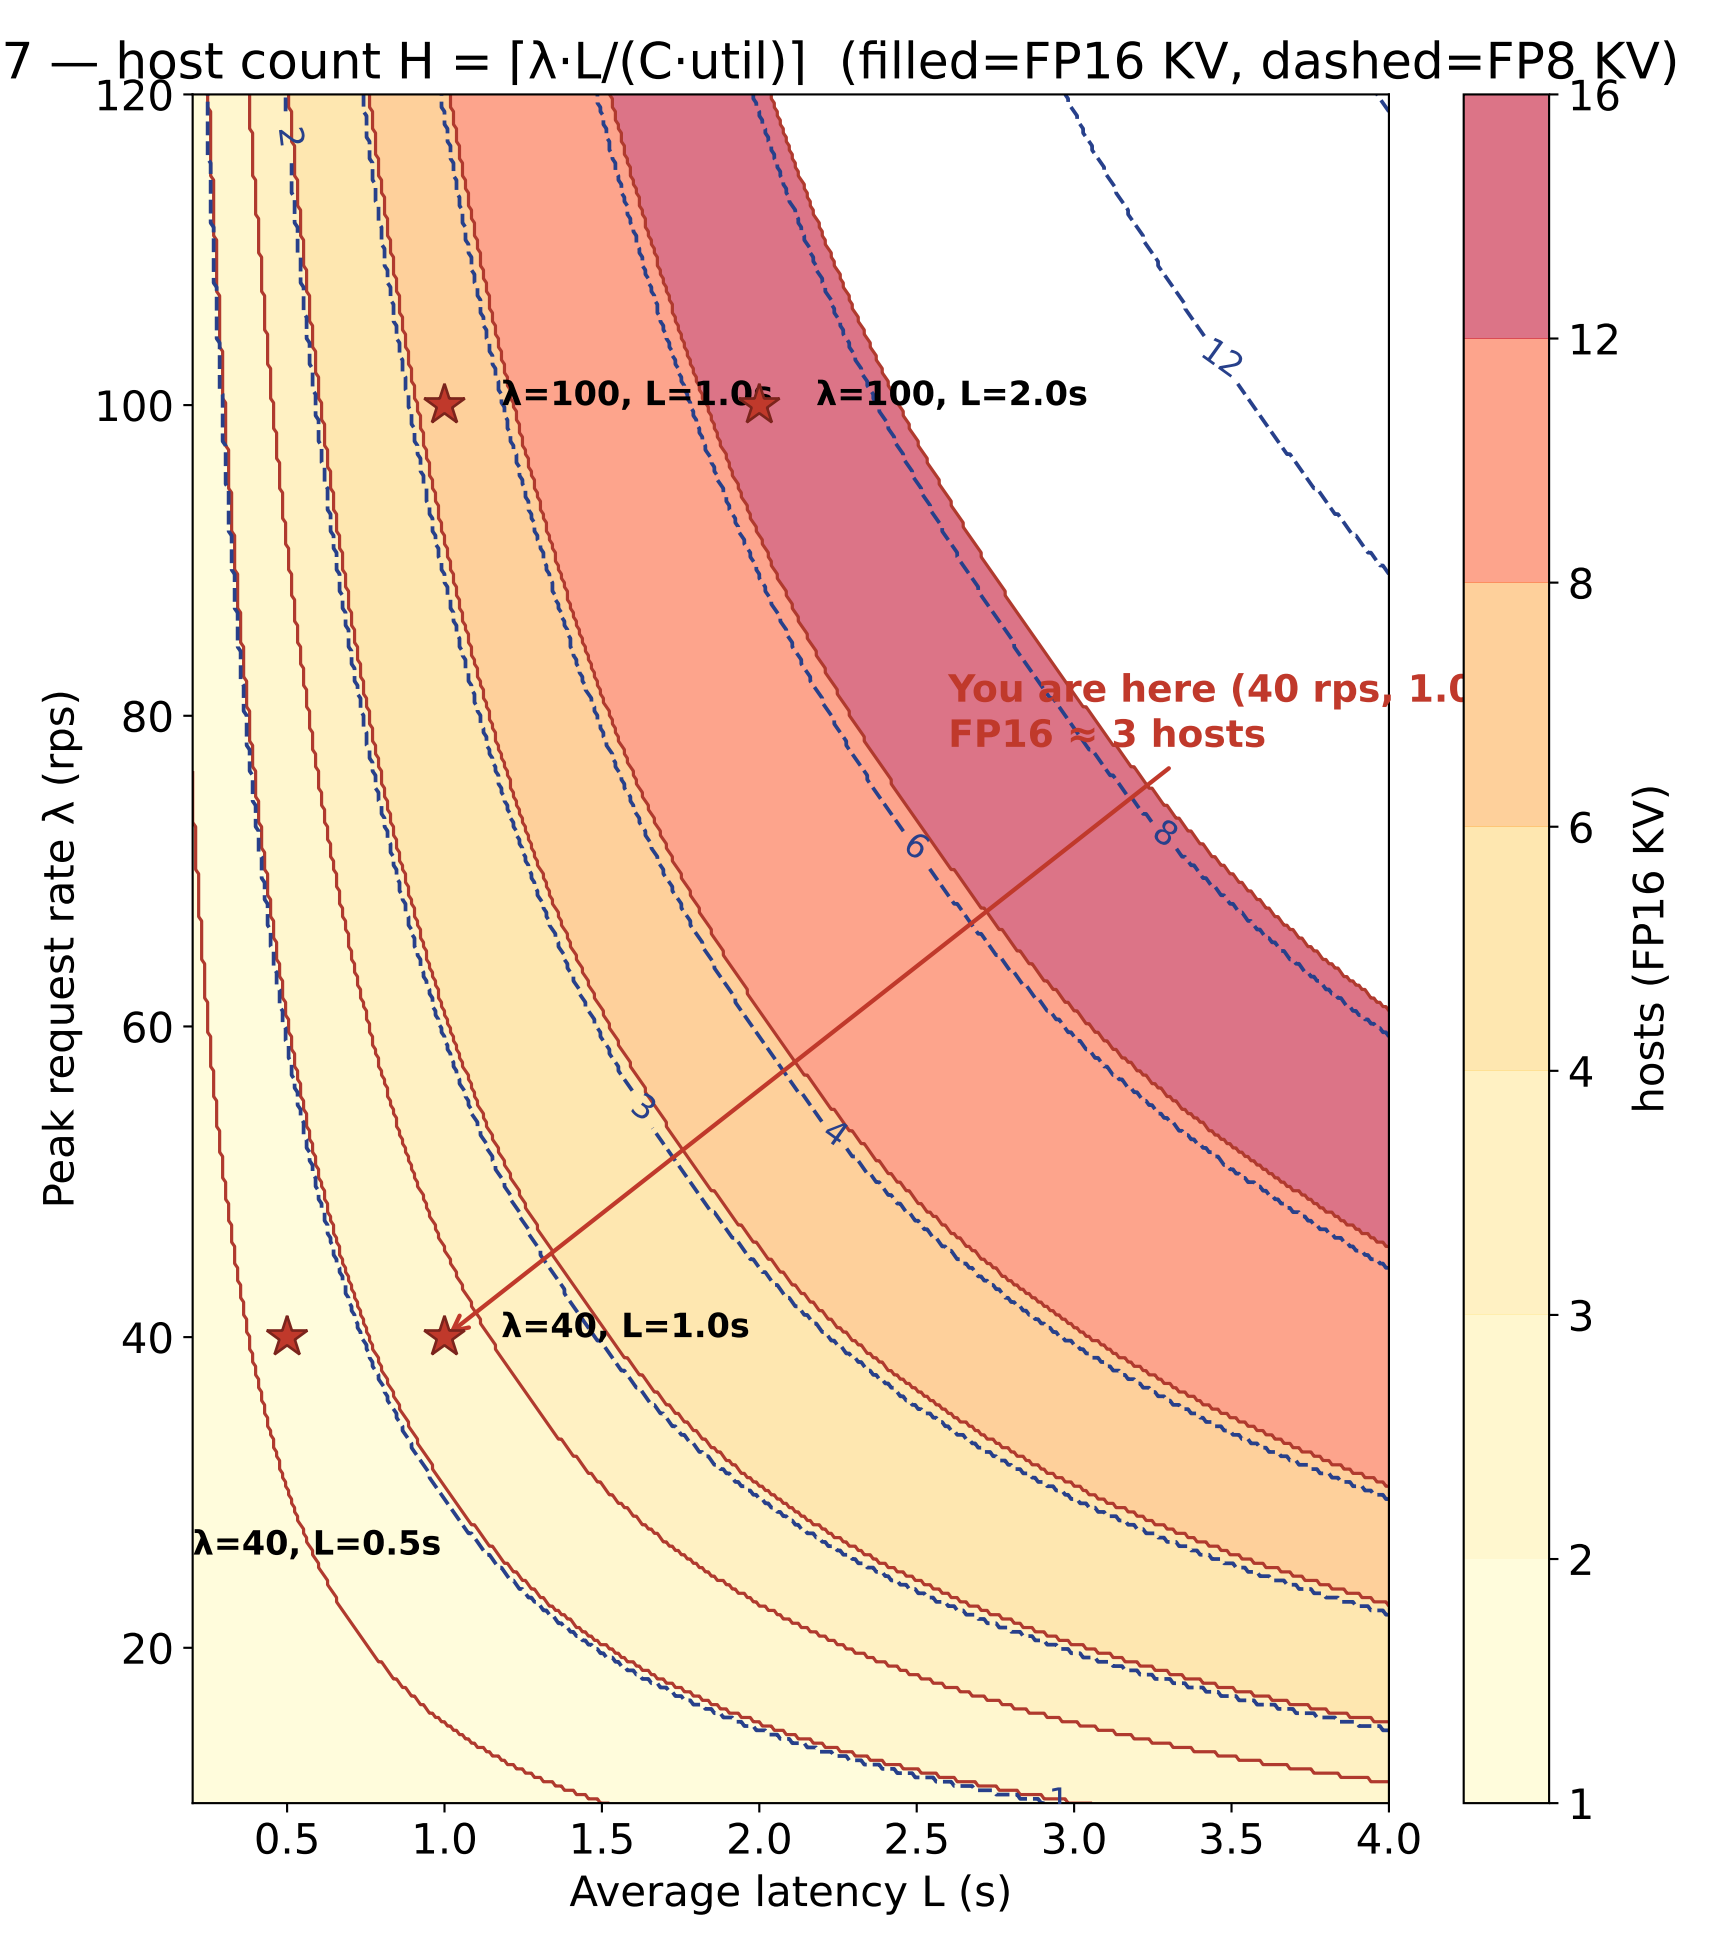

7 — host count H = ⌈λ·L/(C·util)⌉ (filled=FP16 KV, dashed=FP8 KV)
Peak request rate λ (rps)
Average latency L (s)
hosts (FP16 KV)
120
100
80
60
40
20
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
16
12
8
6
4
3
2
1
λ=100, L=1.0s
λ=100, L=2.0s
You are here (40 rps, 1.0
FP16 ≈ 3 hosts
λ=40, L=1.0s
λ=40, L=0.5s