


Fleet capacity (KV/service-time analytical bound, ~ 2.0 req/s/host)



Offered-load utilization (analytical bound, NOT GPU util)

