


Fleet capacity (KV/service-time analytical bound, ~2.0 req/s/host)



Offered-load utilization vs host count (analytical service bound, not GPU util)

