

Full-attention prefill increasingly reflects quadratic attention cost; decode is incurred token-by-token

