

Prefill compute grows super-linearly with context; decode stays flat

