



Prefill per request:

linear + quadratic attention



Decode per token:

fixed 2N (context-independent)



Left panel is PER REQUEST, right panel is PER TOKEN — different units, do not compare directly.

The quadratic term is the textbook $4 \cdot n_{\text{layers}} \cdot L^2 \cdot d$ (matches Chapter 8). [ILLUSTRATIVE][DERIVED]