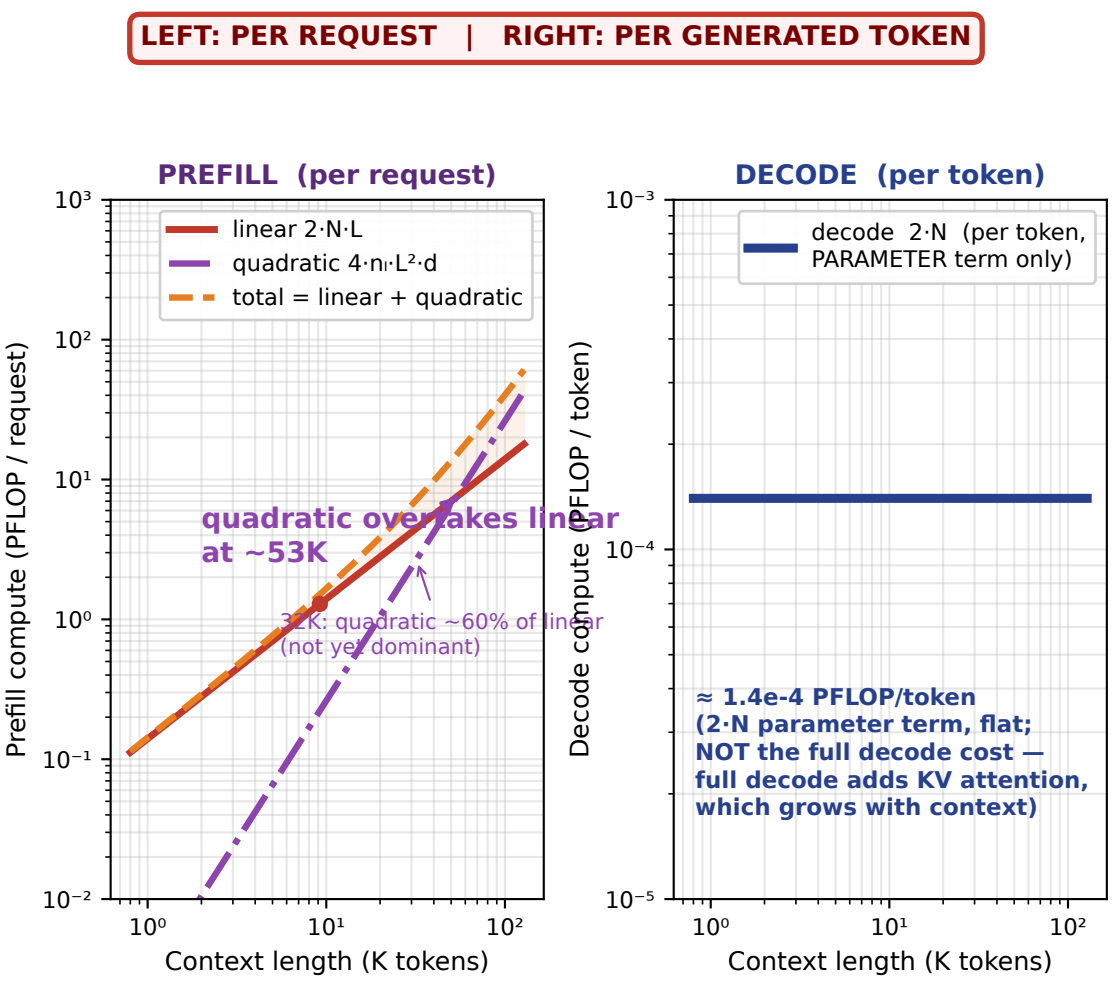
LEFT: PER REQUEST | RIGHT: PER GENERATED TOKEN
PREFILL (per request)
linear 2·N·L
quadratic 4·nₗ·L²·d
total = linear + quadratic
quadratic overtakes linear at ~53K
32K: quadratic ~60% of linear (not yet dominant)
Prefill compute (PFLOP / request)
Context length (K tokens)
DECODE (per token)
decode 2·N (per token, PARAMETER term only)
≈ 1.4e-4 PFLOP/token (2·N parameter term, flat; NOT the full decode cost — full decode adds KV attention, which grows with context)
Decode compute (PFLOP / token)
Context length (K tokens)