Vendor-reported ratios: % remaining

vs each model's own stated reference (predecessor or base model)

NOT A CROSS-MODEL PERFORMANCE RANKING

Each row is relative to that model's OWN stated baseline — no shared scale

One row per reported model — DeepSeek V4's two variants listed separately.

No row for Kimi K3 / Qwen3.8-Flash-Next: no predecessor-relative baseline reported.

reported model	own baseline	KV cache*	Compute*
DeepSeek V4-Flash	*vs V3.2* reported 2026-06	**100% → ~7%** *single-token KV cache*	**100% → ~10%** *single-token inference FLOPs*
DeepSeek V4-Pro	*vs V3.2* reported 2026-06	**100% → ~10%** *single-token KV cache*	**100% → ~27%** *single-token inference FLOPs*
GLM-5.3-Flash	*vs GLM-5.3* reported 2026-08	**100% → ~23%** *BF16 / layer*	**100% → ~33%** *attention only per head / layer*

*Metric scopes differ. 7% remaining = 93% reduction.

Value source: vendor-reported [1P] — DeepSeek V4: arXiv 2606.19348

GLM: z.ai/blog/glm-5.3-flash (2026-08-26). Vendor claims, not reproduced.