Vendor-reported KV / FLOP reduction

Each figure is relative to that model's OWN stated baseline — no shared scale

model family	own baseline	KV / token	FLOP / token
DeepSeek V4-Flash	*vs V3.2*	**100% → ~7%**	**100% → ~10%**
DeepSeek V4-Pro	*vs V3.2*	**100% → ~10%**	**100% → ~27%**
GLM-5.3-Flash	*vs stated ref*	**100% → ~23%**	**100% → ~33%**

Each reduction is relative to the model's OWN stated predecessor — no common denominator.

Value source: vendor-reported [1P] (arXiv 2606.19348; HF zai-org), not yet independently reprofiled.