

Vendor-reported KV / FLOP reduction

Each figure is relative to that model's OWN stated baseline — no shared scale

model family	own baseline	KV / token	FLOP / token
DeepSeek V4-Flash	vs V3.2	100%➡~7%	100%➡~10%
DeepSeek V4-Pro	vs V3.2	100%➡~10%	100%➡~27%
GLM-5.3-Flash	vs stated ref	100%➡~23%	100%➡~33%

Each reduction is relative to the model's OWN stated predecessor — no common denominator.

Value source: vendor-reported [1P] (arXiv 2606.19348; HF zai-org), not yet independently reprofiled.