Corruption is specified; reverse generation is learned

Forward q: add corruption. Reverse p_θ: generate with learned updates.

A Continuous diffusion (schematic state values)

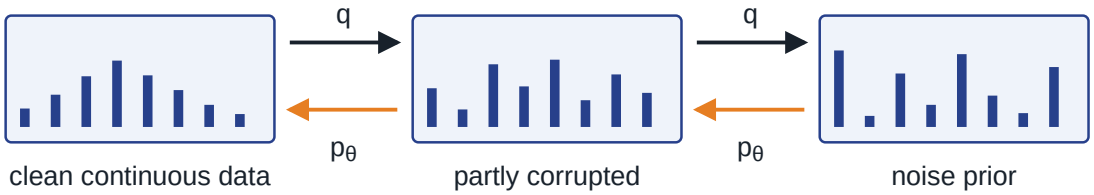

B Masked discrete diffusion (one categorical choice)

M = MASK; text states are illustrative token IDs, not generated results.

Training: construct corrupted inputs; learn reverse predictions.

Sampling: start from the prior; apply learned reverse updates.

Flow matching learns a vector field on a chosen probability path.

Related paths do not make diffusion and flow matching identical.